


Applied Linear Modeling

October 22, 2019

All the things for today

- Mini-lecture on the theory behind the logistic model
- Logistic regression basketball workshop
 - *R* package to download: *odds.n.ends*
- Slides and workshop packet on GitHub



What is logistic regression

- A statistical model used to predict or explain a **binary** outcome variable
- For example:
 - What predicts whether or not someone uses the library?
 - What predicts whether or not someone is a smoker?
 - What predicts whether or not someone owns a gun?

The statistical form of the logistic model

- Because the outcome variable is binary, the linear regression model would not work (it requires a continuous outcome!)
 - Linear model: $y = b_0 + b_1x_1 + b_2x_2 \dots$
- The linear regression model can be transformed using a *logit transformation* based on the logistic function in order to model binary outcomes
- The logistic function is officially defined:

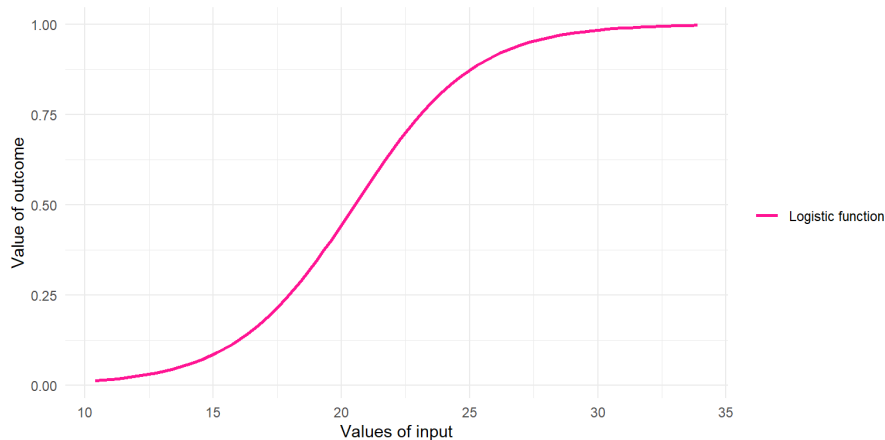
$$\sigma(t) = \frac{e^t}{e^t + 1}$$

- Simplified to:

$$\sigma(t) = \frac{1}{1 + e^{-t}}$$

The logistic function

- The logistic function has a sigmoid shape that stretches from $-\infty$ to ∞ on the x-axis and from 0 to 1 on the y-axis
- The function can take any value along the x-axis and give the corresponding value between 0 and 1 on the y-axis



From the logistic function to the logistic model

- Logistic function to logistic model
 - t is the value along the x-axis of the function
 - $\sigma(t)$ is the value of y for a specific value of t , or the probability of y given t , $p(y)$
 - In the case of logistic regression, the value of t will be the right-hand side of the regression model, which looks something like $\beta_0 + \beta_1 x$
 - Substitute in $p(y)$ and the linear regression model

- Logistic function

$$\sigma(t) = \frac{1}{1 + e^{-t}}$$

- Logistic model

$$p(y) = \frac{1}{1 + e^{-(\beta_0 + \beta_1 x)}}$$

Reading the logistic model

- Logistic model

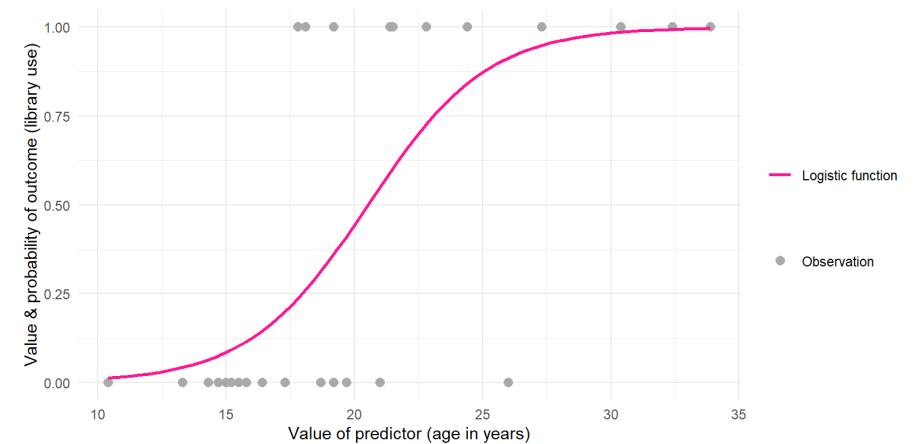
$$p(y) = \frac{1}{1 + e^{-(\beta_0 + \beta_1 x)}}$$

- Where:

- y is the binary outcome variable
- $p(y)$ is the probability of the outcome
- b_0 is the y -intercept
- x_1, x_2 , etc are predictors of the outcome
- b_1, b_2 , etc are the slopes/coefficients for x_1, x_2

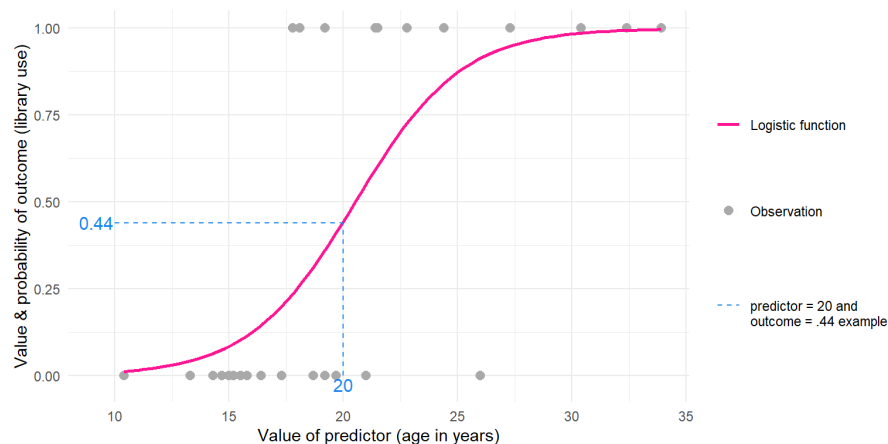
The logistic function with data

- The data points represent library users ($y = 1$) and non-users ($y = 0$)
- The x-axis represents age in years



Example of probability of y for a value of x

- What is the probability of library use for a 20-year old?
 - Starting at 20 on the x-axis, trace a straight line up to the logistic function curve and from there look to the y-axis for a value
 - For the logistic model represented by this graph, the model would predict a probability of y around .44 or 44%



Interpreting the probability

- If this were a model predicting library use from age, it would predict a 44% probability of library use for a 20-year-old
- Since 44% is lower than a 50% probability of the value of y, the model is predicting that the 20-year-old does not have the outcome
- So, if the outcome is library use, the logistic model would predict this 20-year-old was not a library user

Let's play logistic regression basketball!